

This manuscript develops a Visibility Attention Index (VAI), a dictionary-based measure of the share of visible versus functional infrastructure language in Chinese municipal government work reports. The author applies the measure to 6,294 reports from 282 prefecture-level cities between 2002 and 2024. The central claim is that officials emphasize visible infrastructure because such projects are easier for superiors, inspectors, media, or citizens to observe than functional infrastructure such as drainage, heating, water systems, and structural safety.

The core intuition is interesting. Bureaucratic texts may reveal what officials choose to emphasize, and the visible-versus-functional distinction is substantively meaningful. However, the manuscript currently overclaims. It more convincingly measures visible-infrastructure rhetoric than strategic “visibility bias.” The evidence does not yet support the stronger claims about bureaucratic incentives, causal mechanisms, or welfare consequences.

Major Comments

1. Construct Validity

The paper needs to distinguish three claims: VAI measures visible versus functional infrastructure language; this language reflects bureaucratic attention; and this attention reflects strategic response to observability and career incentives. The evidence is strongest for the first claim, weaker for the second, and insufficient for the third.

A city may use more visible-infrastructure language because it actually built more visible projects, because of central urban campaigns, local development needs, fiscal capacity, report templates, or genre conventions. These alternatives do not necessarily imply strategic visibility bias. The paper should therefore describe VAI more cautiously as a measure of visible-infrastructure rhetoric unless stronger evidence is provided.

2. Institutional Premise

The argument depends on government work reports being consequential texts. The manuscript states that they are read by citizens and provincial evaluators and that they communicate performance and commitments. But this is asserted more than demonstrated.

The author should explain who drafts these reports, who reviews them, whether they enter cadre evaluation, whether provincial officials compare them across cities, whether media or citizens actually use them, and whether report content has consequences. The paper should also clarify visibility to whom: superiors, inspectors, citizens, media, or other officials. These audiences imply different mechanisms.

Without this institutional grounding, VAI may simply capture procedural or templated bureaucratic language rather than strategic signaling.

3. Validation evidence useful but weaker than advertised.

Several validation exercises are overinterpreted. The alternative-dictionary correlation is a reliability check, not construct validation. It shows that two similar dictionaries applied to the same genre produce similar scores. It does not show that VAI captures strategic visibility bias. The review-versus-plan differential is interesting but not diagnostic. Retrospective sections naturally describe completed, tangible projects, while prospective sections may use more technical or target-oriented language. This may be a genre difference rather than a behavioral signature. The MOHURD accounting validation is the strongest external check, but the correlation is modest at best. It suggests that VAI contains some information about investment composition, but it does not establish mechanism or direction. The failed Wikipedia validation should be treated more plainly as failed external validation. If Wikipedia was the wrong corpus, it is unclear why it was preregistered as a validation test in the first place.

The manuscript reports human coding in the construction of the lexicon, but it does not appear to provide an independent passage-level validation of the resulting measure. The author should validate the dictionary against a random sample of sentences or paragraphs coded as visible, functional, mixed, or irrelevant, and report precision, recall, false positives, false negatives, intercoder agreement, and examples of ambiguous cases.

4. The additional empirical exercises overextend the paper.

The inspection event study does not support the theory. The narrow TWFE estimate is in the predicted direction, but the result disappears under heterogeneity-robust estimation and reverses sign when the sample is extended. The author reasonably interprets the original result as a staggered-treatment artifact. This should be treated as a failed preregistered test, not as part of the main evidentiary structure.

The CFPS test is also uninformative in its current form. The preregistered idea appears to have been to examine satisfaction with visible versus functional amenities. But the relevant amenity-specific items were unavailable, so the author substitutes broad outcomes such as government evaluation, life satisfaction, and self-rated health. These outcomes do not measure the proposed mechanism.

The welfare calibration should be removed or sharply demoted. It depends on strong assumptions about the social optimum, the welfare function, the mapping from text to investment composition, and the scale of relevant expenditure. The estimates appear more precise than the evidence warrants.

5. The contribution and presentation should be narrowed.

The comparison in Table 0 should be revised. The manuscript implies that prior text-as-data work is largely limited to ideological scaling or generic issue emphasis, but this understates

existing work on expressed agendas, political attention, agenda-setting, bureaucratic responsiveness, audience targeting, and cross-domain policy classification. For example, Grimmer's work on expressed agendas in Senate press releases is explicitly about attention rather than ideology; the broader text-as-data literature has applied automated content analysis to legislative speeches, press releases, media, bills, hearings, treaties, and other political texts; and recent work on cross-domain topic classification explicitly compares priorities across manifestos, speeches, coalition agreements, legislative texts, and social media. The manuscript should therefore frame VAI less as a categorical departure from prior text-as-data instruments and more as a domain-specific measurement application to Chinese municipal government reports.

The paper also relies heavily on preregistration and replication as evidence of rigor. These are valuable, but the PAP, code, lexicons, data documentation, and deviation log must be accessible to reviewers. Yet, none of the links provided in the text actually works or are accessible. This presents a serious reviewability problem because the manuscript repeatedly invokes preregistration and transparent deviations as part of its contribution.

Finally, the manuscript needs substantial editing. It currently tries to be a measurement paper, a formal model, a China bureaucratic politics paper, a validation study, a DiD paper, a CFPS microfoundations paper, and a welfare-calibration paper. These pieces do not cumulate cleanly. The strongest version would define VAI clearly, establish the institutional importance of government work reports, validate the dictionary through human coding and accounting data, and present the review-plan difference as a suggestive descriptive pattern. The failed causal tests, substituted CFPS analysis, and welfare calculation should be moved to an appendix or removed.

In all, I recommend major revision. The paper has a promising measurement idea, but the current version does not yet establish that VAI measures strategic bureaucratic visibility bias rather than visible-infrastructure rhetoric. A publishable version should narrow the contribution, strengthen construct validation, clarify the institutional setting, make preregistration and replication materials accessible, and remove or demote the causal and welfare claims.