

Cover Letter — Journal of Computational Social Science (Revised Submission)

Date: 2026-06-19 **Manuscript:** JCSO-D-26-00240 (Major Revisions) **To:** Editor-in-Chief,
Journal of Computational Social Science (Springer Nature)

Dear Professor Kamihigashi,

Please find enclosed our revised manuscript, “**Visibility Bias in Chinese Municipal Government Work Reports: Measurement and Behavioral Validation**” (formerly *Measuring Visibility Bias in Bureaucratic Text*), submitted for continued consideration as a Research Article. We are grateful to Reviewer #1 for a review that materially reshaped the paper; a point-by-point response accompanies this letter.

What the revision does. The review correctly judged that our original draft measured visible-infrastructure *rhetoric* more convincingly than strategic *visibility bias*, and overclaimed. We have **kept the Visibility Attention Index (VAI) as the paper’s instrument and kept its structure**, and used your comments to **strengthen the validation and narrow the claims** — the paper a reviewer reads is the same one, much better evidenced. Applied to 6,294 Chinese municipal government work reports (282 prefectures, 2002–2024), the contribution remains methodological and, we believe, of direct interest to JCSS’s readership:

1. **A rigorous passage-level validation that reports a precision ceiling.** Against human double-coding, a naive visibility dictionary fails (visible-class precision 0.10, from polysemous appearance words); a concrete, salience-based lexicon improves greatly (0.50 → 0.60) but meets a **polysemy ceiling (~0.60–0.64)** that curation cannot break. Human intercoder agreement is Cohen $\kappa = 0.70$.
2. **A calibrated high-precision alternative.** A three-model **LLM ensemble** attains 0.84 accuracy (inter-model Fleiss $\kappa = 0.835$), benchmarked against the same human gold standard — a concrete precision-vs-transparency trade-off.
3. **Behavioral (criterion) validity as the decisive test.** Under quasi-exogenous, retirement-driven secretary turnover, the *valid* concrete text measure co-moves with *real* accounting-based cosmetic investment (both up the year after turnover, clean pre-trends), while the *naive* measure does not respond ($p \approx 0.5$). We argue, and show, that behavioral co-movement — not passage precision or internal reliability — is the decisive validity test for measures of strategic attention.
4. **Honest scope and reproducibility.** Failed and assumption-dependent analyses (an inspection event study, a CFPS micro-foundation, a structural welfare calibration) are demoted to appendices and reported transparently; a large-scale procurement check (Zhejiang public bidding, ~2.96M records) corroborates the compositional tilt on both project frequency ($\approx 2:1$) and investment amount.

Author team. Since the original submission we have added a coauthor whose work answers needs the review surfaced: **Liu Can** (co-first; urban & rural planning) built and verified the

procurement-classification scheme (Section 4.3) and provided the independent second coding behind the intercoder agreement. Contributions follow ICMJE and are listed in the Statements.

Originality and concurrent submission. The work is original and not under consideration elsewhere. A pre-registered predecessor (OSF ZMJY5) and an earlier urban-economics-framed draft (desk-rejected at *Habitat International*, 2026-04-13, on scope grounds) are disclosed in Appendix D; the present paper is a substantial methodological repositioning, not a resubmission of that draft.

Reproducibility. Pre-registration: OSF DOI [10.17605/OSF.IO/ZMJY5](https://doi.org/10.17605/OSF.IO/ZMJY5). Replication archive (instrument, gold-standard codings, all analysis and re-pull code, procurement classification, deviation log): Zenodo concept DOI [10.5281/zenodo.19569978](https://doi.org/10.5281/zenodo.19569978) (always resolving to the latest version), mirrored at GitHub ([andyhsi2023-cq/visibility-bias-v2](https://github.com/andyhsi2023-cq/visibility-bias-v2)). We have re-deposited and verified that every link resolves: the Zenodo record is now **open access** and the GitHub mirror is **live and public**. We apologize that the original record was not accessible to the reviewer and have corrected it.

Use of generative AI. Large-language-model classification is an *object of study and a measurement tool* in this paper (the §3.2.0 ensemble), documented and reproducible; separately, the authors used an LLM for copy-editing only. All substantive content and conclusions are the authors' own.

We suggest, as before, Daniel Mattingly (Yale), Jennifer Pan (Stanford), and Yuhua Wang (Harvard) as reviewers, and note any recusals in `suggested_reviewers.md`. Thank you for the opportunity to revise.

Sincerely,

Hongyang Xi (corresponding author), on behalf of both authors — Hongyang Xi and Liu Can (co-first authors) Chongqing Survey Institute Co., Ltd., Chongqing, China Email: 26708155@alu.cqu.edu.cn · ORCID: [0009-0007-6911-2309](https://orcid.org/0009-0007-6911-2309)